

# Fork() System Call and Processes multi-taskings

# Programs and Processes

- Program: Executable binary (code and static data)
- Process: A program loaded into memory
  - Program (executable binary with data and text section)
  - Execution state (heap, stack, and processor registers)

## Program

```
int foo() {  
    return 0;  
}  
  
int main() {  
    foo();  
    return 0;  
}
```

## Process

```
int foo() {  
    return 0;  
}  
  
int main() {  
    foo();  
    return 0;  
}
```

Heap

Stack

Registers

# fork()

- A system call that creates a new process identical to the calling one
  - Makes a copy of text, data, stack, and heap
  - Starts executing on that new copy
- Uses of fork()
  - To create a parallel program with multiple processes (E.g. Web server forks a process on each HTTP request)
  - To launch a new program using exec() family of functions (E.g. Linux shell forks an 'ls' process)

# fork() example

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    int seq = 0;
    if(fork()==0)
    {
        printf("Child! Seq=%d\n", ++seq);
    }
    else
    {
        printf("Parent! Seq=%d\n", ++seq);
    }
    printf("Both! Seq=%d\n", ++seq);
    return 0;
}
```

```
>> ./a.out
Parent! Seq=1
Both! Seq=2
Child! Seq=1
Both! Seq=2
```

- Differentiate child and parent using return value of fork()
  - “Child” process return value is 0
  - “Parent” process gets child’s process id number
- **Copies** execution state (not shares)
  - Child copies stack variable seq onto its own stack
  - Child / parent has own copy of seq

# How do Parent / Child Run in Parallel?

- Obvious answer: by running on different processors on a multiprocessor system
- What if it's a uniprocessor system?
- What if other processors are already busy?
- Answer: by context switching (running each process in short time slots)

# Spawning a New Program

- Combination of `fork()` and `exec(...)`
  - `fork()`: Clone current process
  - `exec(...)`: copy new program on top of current process
- `Exec(...)` family of functions
  - `execl`, `execvp`, `execle`, `execv`, `execve`, `execvp`
  - User space wrappers for the **execve** system call
  - Also called the “program loader”
    - Loads in text and data sections of a binary executable
    - Links in any shared objects and perform relocations
    - Sets up stack and starts executing
  - What Linux shell calls when launching a program

# execvp() example

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    if(fork()==0)
    {
        char *args[3] = {"ls", "-al", NULL};
        execvp(args[0], args);
        // DOES NOT GET HERE
    }
    else
    {
        printf("Parent!\n");
    }
    printf("Only parent!\n");
    return 0;
}
```

```
>> ./a.out
```

```
Parent!
```

```
Only parent!
```

```
drwx----- 4 wahn UNKNOWN1  4096
Oct 21 08:13 .
```

```
drwxr-xr-x 10 wahn UNKNOWN1  2048
Oct 21 08:13 ..
```

```
-rwxr-xr-x 1 wahn UNKNOWN1   6743
Oct 21 08:12 a.out
```


- execvp never returns since memory is overwritten using another program image

# Wait() function to synchronize parent -> child termination

A call to wait() blocks the calling process until one of its child processes exits. After child process terminates, parent **continues** its execution after wait system call instruction.

Child process may terminate due to any of these:

- It calls exit()
- It receives a terminating signal (from the OS or another process)
- It returns (an int) from main



## Syntax in c language:

#include #include // take one argument status and  
returns // a process ID of dead children.

```
pid_t wait(int *stat_loc);
```